

Swapping AttnGrounder's Attention for a Brain-Inspired Associative Memory (BDH)

Testing the “Grounding = In-Context Retrieval” Hypothesis on Talk2Car

Sudarshan Sudhakar

July 20, 2026

Outline

- 1 Motivation & Hypothesis
- 2 Method
- 3 How We Measure Success
- 4 Results
- 5 Where We Stand vs. SOTA
- 6 What's Novel Here
- 7 Methodology
- 8 Future Scope

The Task: Visual Grounding

Visual grounding (Talk2Car): given an image and a natural-language command, localize the single object the command refers to.

“Turn left behind the silver car” → one bounding box.

The hard cases are scenes with **multiple objects of the same class** — the model must disambiguate using language, not just detect.

Baseline model

AttnGrounder (Mittal, IIT Mandi) — single-stage, Darknet-53 + BiLSTM/GloVe, a lightweight *visual-text attention* module fuses the two modalities before a YOLO head.

Why Bring in BDH?

Dragon Hatchling (BDH) — Kosowski et al., Pathway 2025


A brain-inspired sequence architecture whose working memory is an **explicit, growing outer-product associative memory** ($\rho = K^\top V$), updated with Hebbian-style read/write — rather than a fixed-size compressed state (e.g. Mamba/SSMs).

The hypothesis we're testing

“Grounding = in-context retrieval”: architectures with explicit associative memory should have an edge precisely when a query must be matched against **multiple similar candidates** in context — because that is a retrieval problem, not a compression problem.

Prediction: swapping AttnGrounder’s attention for a BDH module should raise AP50 *disproportionately* on scenes with several same-class candidate objects (“ambiguous” scenes), relative to scenes with only one.

Architecture: A Drop-In Swap



Baseline attention: softmax cross-attention (region queries $\times$ word keys), scalar gate β .

BDH module: symmetric or causal-context kernel (3 modes tried, see Results); ReLU-sparse addressing; explicit outer-product memory; multiplicative bilinear gate. Same output shapes $\Rightarrow$ true drop-in.

Keeping the Comparison Fair

- **Single component swap:** backbone, text encoder, YOLO head, losses, hyperparameters — all byte-for-byte identical between variants.
- **Fusion trim:** BDH fuses only $[f_{visu}, f_{lang_attn}]$ (2 streams) instead of the original 3, to offset most of the module's added parameters.
- **Same training recipe for all variants:** A100-40GB, bf16, 100 epochs, Adam (lr 10^{-4} , backbone lr/10), batch 14, poly LR schedule.

Variant	Params	vs. baseline
Baseline (AttnGrounder, unmodified)	75.84M	—
BDH (any mode, single-head)	76.63M	+0.79M (+1.0%)


Backbone dominates parameter count either way — this is a fusion-mechanism comparison, not a capacity comparison.

The Metrics, Explained

AP50 (our headline number)

A prediction counts as **correct** if the predicted box overlaps the ground-truth box by

Intersection-over-Union ≥ 0.5 . AP50 = % of commands localized correctly.



Params

Total trainable parameters — proxy for model size / compute cost. Lower is “lighter.”

Ambiguous vs. Unambiguous (Step 6)

The key stratification for testing our hypothesis —
ambiguous: ≥ 1 other same-class object in frame;
unambiguous: referred object is the only one of its class. Built from nuScenes 3D metadata — see next slide.

Step 6: Stratifying the Val Set

Talk2Car's own annotations only record the *referred* object — not what else is in the scene. To test the hypothesis we need to know, per image, how many **other** objects share the referred object's class.

Data pipeline built for this


Join Talk2Car commands → nuScenes v1.0-trainval metadata (front-camera 3D annotations) via the sample token → count same-class objects visible in CAM_FRONT.

Bucket	Definition	Count (val, $n=1163$)
Ambiguous	≥ 1 other object, same class	830
Unambiguous	referred object is unique in class	333

Ablation Study — Full Results

Three BDH kernel formulations trained and evaluated end-to-end (same recipe, same data, same eval pipeline as baseline):

Variant	AP50 (all) <i>n</i> =1163	AP50 (ambig.) <i>n</i> =830	AP50 (unambig.) <i>n</i> =333	vs. base (all)
Baseline	64.92	63.37	68.77	—
BDH Mode C	64.06	62.17	68.77	−0.86
BDH Mode A	65.09	63.98	67.87	+0.17
BDH Mode B	60.53	58.19	66.37	−4.39



Bar chart: headline pair (baseline vs. winning Mode A). Full 3-mode sweep in the table above.

Interpretation

- **Mode A** is the first (and only) formulation pointing in the hypothesis's predicted direction: beats baseline on ambiguous scenes (+0.61) while ceding ground on unambiguous (−0.90) — a genuine disambiguation trade-off.
- **Modes B and C both do worse specifically on ambiguous scenes** (−1.20 and −5.18 vs. baseline) *despite both having* an explicit outer-product associative memory, same as Mode A.

Key insight

The win isn't just “having associative memory” — it's specifically Mode A's **causal text-self-contextualization** step (words attend to each other before regions query them). This directly motivates the BDH→attention hybrid idea (Future Work).

Caveat

Single seed per variant — a ~ 1 AP50 gap on 830 images carries meaningful noise. Encouraging, not yet conclusive.

Where We Stand — Same-Class Comparison

ThinkDeeper (Liao et al., Nov 2025, ViT+BERT+world-model+hypergraph, **76.64** AP50) is #1 on Talk2Car — but its own Table 1 shows a clear bracket structure by backbone class:

Model	Backbone	AP50 (test)	Params
AttnGrounder (paper)	CNN	61.32	75.84M
Ours: BDH Mode A	CNN	65.09 (val)	76.63M
CMSVG	CNN (EfficientNet)	68.61	—
VL-BERT	CNN+BERT	70.03	~118M
RSD-LXMERT	LXMERT	72.64	~210-230M
ThinkDeeper	ViT+BERT+...	76.64	not stated

Reframed target

Not “reach 76.64” (needs a full transformer/VLM backbone swap — a different project). Realistic, honest target: **beat CMSVG’s 68.61** — top of our same-class (CNN, non-VLM) bracket.

We’re also meaningfully smaller than the transformer-class models (60-70% of VL-BERT, ~1/3 of RSD-LXMERT) and trivially smaller than VLM-based ones (7B–72B params).

What's Genuinely New Here

Based on our literature review — framed as “no prior work found,” not an unfalsifiable “first ever.”

Headline novelty 1: BDH applied to cross-modal grounding

BDH's own authors tested it on **single-modality** vision classification (E-BDH, image-only) and language modeling — never on cross-modal vision-language fusion. No evidence found of BDH applied to visual grounding, or to Talk2Car, anywhere in the literature. This is the first cross-modal grounding application.

Headline novelty 2: three original non-causal kernel adaptations

BDH was designed causal/autoregressive (language has sequence order; grounding doesn't). Modes A/B/C are original designs adapting BDH's associative-memory mechanism to a non-causal, bimodal setting — not a plug-and-play reuse.

- **Significant — quantitative ambiguity stratification:** nuScenes-joined same-class-object counting on Talk2Car's own val split (not shipped with the dataset) — enables testing the hypothesis at the sub-population level, not just in aggregate.
- **Significant — the fusion trim as a confound control:** deliberately holds params near-parity (+1%) with baseline, so the mechanism swap is the *only* variable — unlike most SOTA papers, which change backbone + mechanism + data together, making it hard to attribute their gains to any one factor.

A Note on Methodology: Why Val, Not Test

- Talk2Car's test-set ground truth has **never been publicly released** — the only historical evaluation path was an Alcrowd-hosted server (ECCV 2020 C4AV challenge).
- That server **closed to submissions in August 2020**. No evidence of a live replacement.
- Papers reporting Talk2Car “test” scores since then (RSD-LXMERT 2022, CAVG/ThinkDeeper 2024–2025) still point to that same dead link — their numbers are most likely **self-reported, not independently verified**.

Our choice

Report **val-set AP50 only**, explicitly caveated against other papers' self-reported test numbers. Less flattering, more honest — directionally informative, not strict apples-to-apples.

In progress

Tversky + Focal loss (Salehi et al. 2017 / Lin et al. 2017 — ThinkDeeper's own aux-loss references) replacing plain BCE on the β /attention-map supervision. One loss swap, no architecture change — targets foreground/distractor imbalance, plausibly most relevant in ambiguous scenes. Training now on Mode A.

- **BDH stacking** (2-3 layers): current module is a single lift→read→gate→decode pass. Tests “does more memory depth help.”
- **BDH→attention hybrid**: motivated by today's insight — memory builds text context, then standard attention aligns.
- **Depth/distance prior**: reuse the 3D data already pulled from nuScenes (Step 6) — no new model needed.
- **Pretrained text encoder** (DistilBERT/MiniLM): replaces 2014-era GloVe+BiLSTM — contained scope, doesn't touch BDH, vision backbone, or detection head.

Scoped out for now: YOLOv7/YOLO26 backbone+head swap — legitimate lever, but a much bigger, separate engineering track that doesn't test the memory hypothesis.

Why Keep Pursuing This

- ① **We already have a real, encouraging signal.** Mode A shows the exact hypothesis-predicted trade-off — not just infrastructure readiness, an actual result worth building on.
- ② **The infrastructure cost is already paid.** Data pipeline, drop-in BDH module (3 modes), training loop, and the stratified evaluator are all built and validated end-to-end. Remaining work is experiment iteration, not engineering.
- ③ **Low marginal model cost.** BDH adds $\sim 1\%$ parameters — stays a clean, fair test of the *mechanism*.
- ④ **Directly tests a live hypothesis** from recent literature (BDH / Dragon Hatchling, <1 year old) in a task setting the original work never evaluated.

Summary

- Built a validated, drop-in BDH swap (3 kernel modes) for AttnGrounder, plus a from-scratch baseline reproduction (64.92 AP50 vs. paper's 63.30%).
- Built a nuScenes-based pipeline to stratify Talk2Car val by scene ambiguity (830 ambiguous / 333 unambiguous).
- **Mode A wins the ablation:** beats baseline on ambiguous scenes specifically (+0.61) — the first real signal supporting the hypothesis. Modes B/C do not.
- We're competitive with the CNN-class SOTA bracket (61–69 AP50) at a smaller parameter count than most peers.
- Next: Tversky+Focal loss (running now), then BDH stacking, BDH→attention hybrid, and a pretrained text encoder.

Thank you — questions welcome.